AI DEVELOPMENT FOR PHOTOVOLTAIC INSTALLATION FORECASTING

Guillermo Blanco[1][0009−0003−1013−6892], Marco Antonio Melgarejo[2,3][1111−2222−3333−4444], and Third Author[3][2222−−3333−4444−5555]

[1] Princeton University, Princeton NJ 08544, USA
[2] Springer Heidelberg, Tiergartenstr. 17, 69121 Heidelberg, Germany
`lncs@springer.com`
http://www.springer.com/gp/computer-science/lncs
[3] ABC Institute, Rupert-Karls-University Heidelberg, Heidelberg, Germany
`{abc,lncs}@uni-heidelberg.de`

Abstract. The photovoltaic industry is one of the mains inside the field of renewable energy, with Spain being one of its headliners. Therefore, the planification of the photovoltaic plant production is key to administrate the energetic plan of the beneficiary.
This project will detail the development of a prediction model, trained through the exploit of data coming from pilot photovoltaic plants, floating or fixed, both real and belonging to CTC (Centro Tecnológico de Componentes) pilot plants. In base of this model, it will be possible to know the expected generated power and administrate the energetic charges in an optimum way.
The content will focus both on the tools used in the development and the project's life cycle, exposing the difficulties found through the process and how they were solved.

Keywords: Machine Learning · Photovoltaic · Artificial Intelligence.

1 Introduction

The spanish photovoltaic market is divided between productors and self-consumers. Productors are those who sell their energy to the electrical network. On the other hand, self-consumers do not sell their production, instead, they save it for themselves. In Spain, a regulatory barrier has been implemented, restricting photovoltaic installations. Some self-consumers are forced to comply with the zero-injection requirement, ensuring no power is fed into the grid.

The zero-injection requirement is a law which prohibits the injection of the production into the electrical network. Therefore, the self-consumers must adjust the photovoltaic production in order not to exceed the electrical demand of the installation [1].

The energetical flexibility in this context defines the capacity of adaptation of the demand depending on external agents such as rates or impositions. It has been proved that it is possible to save around 15-25% of the total spend when a predictive control of the production is being used, while it also contributes to lower the CO_2 emission [8].

Optimization is divided in two types, according to its duration. The short-term optimization focus on hourly energetical production to minimize costs and increase efficiency in a short period of time. The long-term optimization tries to determine the optimal dimensions of the installation and components while minimizing the cost of installation. In order to calculate the electrical demand of an installation, we need to consider multiple factors depending on time:

$$D(t) = B(t) + P(t) + N(t) + A \tag{1}$$

Where B(t) represents the basic work cycle, P(t) the peak demands, which introduce the maximum demand events without overlapping, N(t) the noise, which measures the stochastic noise fluctuations. Last, A means the base consumption, assuring realistic conditions [2].

Spain has gone from 11GW of photovoltaic installed power in 2020 to more than 32GW by the end of 2024 [4], becoming the main production technology, even surpassing the eolic. Nowadays, photovoltaic energy means the 17% of the total national energy, making Spain one of the headliners in solar energy in Europe.

Internationally, photovoltaic represents 46% of renewable energy,and the IEA (International Energy Agency) estimates around 3.6TW of new renewable energy before 2030, being 80% of this photovoltaic. As 2030 gets closer, photovoltaic industry will increase its presence as the main renewable type of energy.

Comparing to other renewable energies, photovoltaic gets much more positive results in environmental impact. Polluting during its manufacturing process are lower than other types of energy, such as eolic. In 2019, it was believed that photovoltaic had prevented 7 million of tons of CO_2 to be casted to the atmosphere, amount equivalent to 3 million cars circulating throughout a year [10].

In spite of having accomplished this achievements, photovoltaic industry still has some challenges, being the prize of the energy one of the main. Being a tecnological sector in such a big proliferation, internal competence has raised significantly, which has caused a fall in the sell prize of solar energy. This situation has directly affected the installations profitability [9]. If harsh actions are not taken soon, the situation is expected to worsen by the end of 2026.

The goal of this project is to create a prediction model that helps installations under the zero inject rule to administrate their demand through forecasts. It must predict the expected value of the power generated by the solar installation.

2 Related work

Photovoltaic power forecasting has gained prominence amid Spain's rapid solar expansion. Early models employed statistical data but limited by non-linear weather dependencies.

Machine learning advancements shifted toward tree-based ensembles. Random forests and Extra Trees improved handling of meteorological inputs like irradiance and temperature, achieving robust short-term predictions on UCI datasets. Gradient boosting, particularly XGBoost, outperformed these via optimized regularization and feature importance, with R^2 scores exceeding 0.90 in benchmarks [11].

In Spain, tools like Shirenda_PV leverage ML on REE data (2015-2020) for nationwide yield estimation, aiding zero-injection compliance under RD 244/2019. Recent works forecast rooftop PV with tuned XGBoost, emphasizing real-time AEMET integration, yet overlook edge deployment [12].

These contributions enable demand optimization, saving 15-25% costs, but gaps persist in low-resource hardware forecasting and hybrid skforecast with XGBoost for real pilot data, motivating our methodology.

Table 1. Comparison of related forecasting studies

Work	**Method**	**Horizon**	**Key Metric**	**Limitation**
Chen (2016) [12]	XGBoost	Short-term	R^2: 0.92	Offline focus
Zhang (2023) [13]	LSTM-AdaBoost	Short	RMSE: 0.043	Synthetic data
Shirenda_PV (2025) [14]	ML on REE	Long-term	N/A (yield est.)	No real-time

3 Methodology

3.1 Preprocessing

When trying to create a prediction model, the first step to take is to make sure that the dataset which will be used to train it is prepared and processed. This must have similar conditions to the place where we try to predict, in order for the model to create patterns of that place in particular.

Once we have the final processed dataset, an analysis is needed, in order to find and eliminate possible outliers which would harm the precision of the model. The boolean formula to detect outliers is:

$$outlier? = (value < (percentile25 - 1.5 * IQR)) | (value > (percentile75 + 1.5 * IQR)) \tag{2}$$

Percentile25 is the value below which is the 25% of the data. Percentile75 is the same, with 75% instead. The IQR (interquartile range) is the difference between them both [3]. Removing the outliers causes a not too large reduction of the dataset, which is favorable.

There are some samples with a value of 0 that go against logic, probably due to bugs in the register system. We need to remove them, because they would increase the prediction error falsely. The way used to correct them was through linear interpolation, according to the following formula [5]:

$$f(x) = f(x_0) + \frac{x - x_0}{x_1 - x_0} * (f(x_1) - f(x_0)) \tag{3}$$

Here is an example of the result:

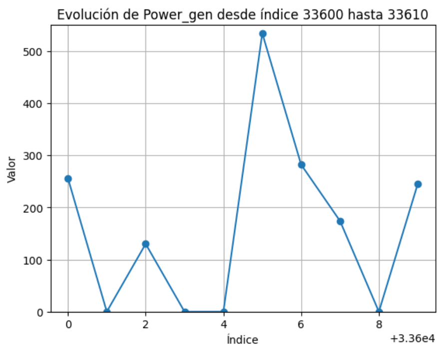

Fig. 1. Before interpolation

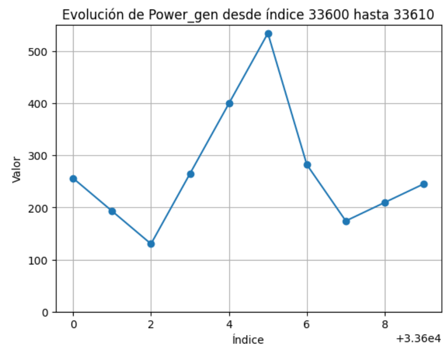

Fig. 2. After interpolation

The data was obtained from a real photovoltaic installation on the roof of CTC, throughout a whole year. It was necessary to put together the dataset of power generated and the meteorological dataset, and sincronize them temporarily. This was repeated for another photovoltaic installation in Lorca, but on this paper we will focus on the one in Santander. Finally, the dataset had some of its rows with a NaN value in one or more columns, so the final step was to remove the rows where this happened. After having done all the processing, this is the final dataset we obtained:

Table 2. Dataset descriptive statistics

Statistic	Power_gen	temperature	windSpeed	solarRad	humedad	rain	pressure
count	41821.000000	41821.000000	41821.000000	41821.000000	41821.000000	41821.000000	41821.000000
mean	86.632242	16.318428	6.869420	101.447096	79.292532	0.525492	30.130212
std	156.516478	4.222477	5.991205	162.073291	10.709026	1.244369	0.196403
min	0.000000	4.440000	0.000000	0.000000	46.000000	0.000000	29.600000
25%	0.000000	12.900000	3.220000	0.000000	73.000000	0.000000	30.000000
50%	0.000000	16.300000	4.830000	0.000000	82.000000	0.000000	30.200000
75%	94.000000	19.800000	9.660000	153.000000	88.000000	0.000000	30.300000
max	780.000000	30.600000	27.400000	631.000000	96.000000	5.400000	30.700000

Due to the fact that the variables of the dataset are scaled differently, is a good practice to scale them. Its crucial to perform this after spliting into train and test data, because, if doing so before, we would be leaking information from the training data into the test data. This practice is known as data leaking, and it decreases the performance of the model.

Scaling makes all variables go within the same range, commonly between 0 and 1. To achieve this we will use the MinMaxScaler, whose formula is [7]:

$$x_{\text{iesc}} = \frac{x_i - x_{\text{imin}}}{x_{\text{imax}} - x_{\text{imin}}} \tag{4}$$

After the processing, we can begin the developing and training process of the model.

3.2 Developing

Temporal series are data ordered chronologically. Forecasting consists on predicting the future value of a temporal serie, being able to use its past behaviour autoregresively and external variables [6]. We will use exogenous variables such as solar irradiance, temperature or wind speed, as well as endogenous variables, like the mean, median, maximum or minimum value. Temporal series differ from classic regression model because of the fact that the pass of time plays a starring role, but, in the end, they are regression models, like Gradient Boosting or Random Forest.

A forecaster model was trained using a XGboost. In order to calculate the optimum parameters of the XGboost a GridSearchCV was used. After the training, we could see the importance of each variable and which one contributed more to the result of the model:

Table 3. Feature importance

Variable	Importance
Lag_1	0.639488
Lag_2	0.158001
solarRad	0.054222
Hour	0.032133
Lag_3	0.013434
Month	0.012951
Roll_min_20	0.009755
Lag_6	0.008500
Lag_7	0.008080
Roll_max_20	0.007707
Other variables	0.055729

It is clear that the most important variables are the samples which precedes the predicted value, which is normal, due to the fact that this variables have the value of the objective variable 5 and 10 minutes earlier. The rest of the variables have little importance, being the solar irradiance and the hour of the day the ones following the lags. This makes sense physically. The minimum and maximum value of the 20 previous samples makes it also into the top 10. Although it would seem that we could only use the Lag_1 and Lag_2 variables, this would be an error because, even though the rest of the variables have little impact on the result, they help to reduce the error in prediction.

It is vital to make sure that we are not overfitting our model. This means not making the model memorize the training data instead of learning prediction patterns. This would result in great metrics when using the same data to test, but a complete failure once we test the model with extern data. One way to prove the non-existance of overfitting is observing how the error in validation and training data progresses through iterations. If the error increases considerably in the validation but continues to fall in the training, we have an overfitted model. We can see that we are not in front of an overfitted model:

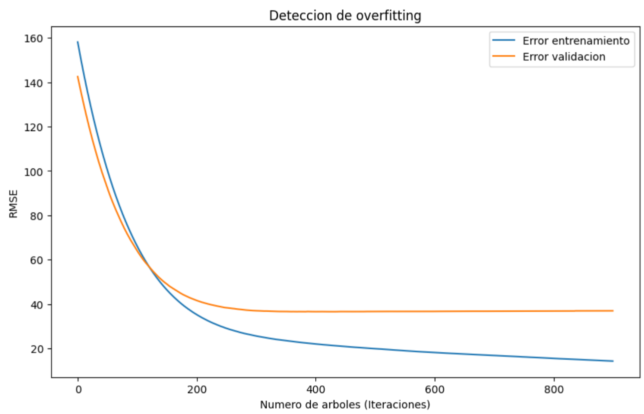

Fig. 3. Training and validation error across iterations.

4 Results

4.1 Prediction metrics

After the prediction of the test data, we got the next results:

Table 4. Prediction metrics

MAE	16.5215
MSE	1631.2590
RMSE	40.3888
R^2	0.9212

The results are not perfect but prove that the model is robust and covers a good part of the variable variance.

4.2 Graphical representation

A good way to prove the accurracy of the model is to put next to each other the real values and the predicted ones. The figure below represents in red the predicted values, and in blue the actual value of the test data.In general, this was the whole test data:

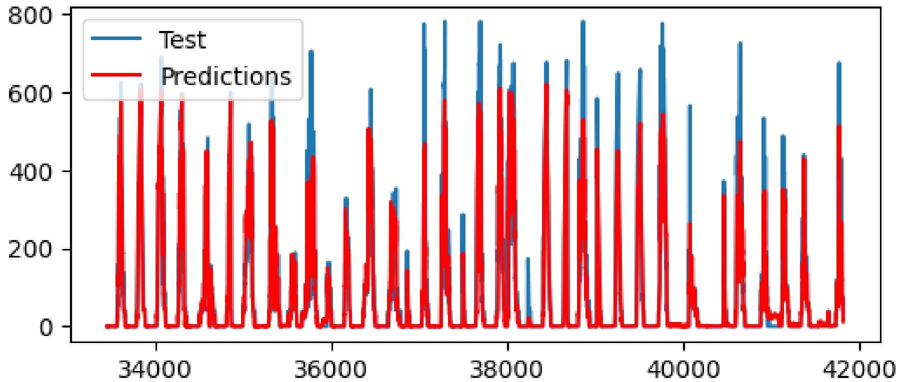

Fig. 4. Prediction through the whole dataset.

It is interesting to prove the accurracy at a lower level, trying to predict the objetive variable in the next 24 or 48 hours. This works as the figure before, red for prediction and blue for real value. Below are the predictions for 24 and 48 hours ahead, what could be the most useful information that the installation may need. We can see that they have a very similar form, proving that the model is capable of predicting the tendency of the power generated by the installation:

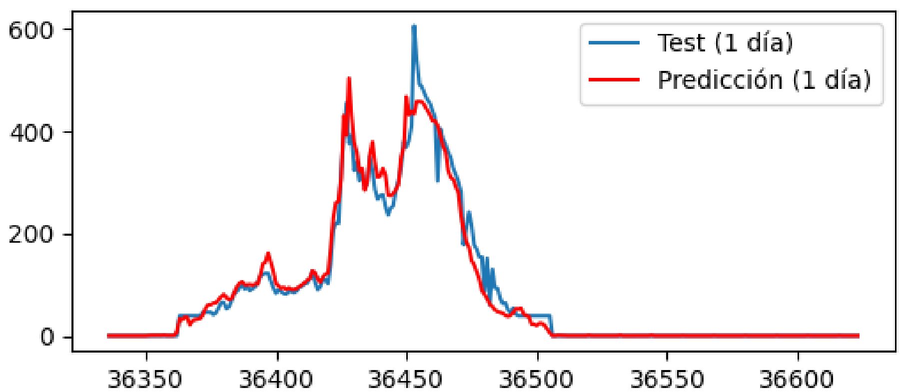

Fig. 5. Prediction of the next 24 hours.

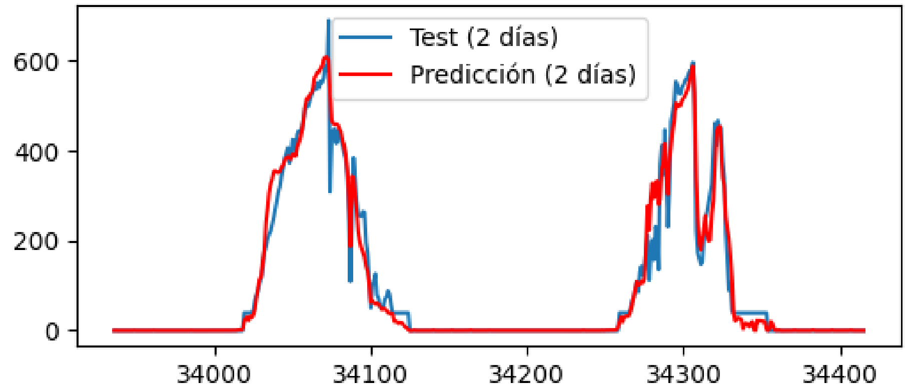

Fig. 6. Prediction of the next 48 hours.

Another way to visualize how the prediction went is not through time but putting all the data together. The more it looks like the discontinuous red line, the better. On the axis Y is the predicted value, and on the axis X the real value:

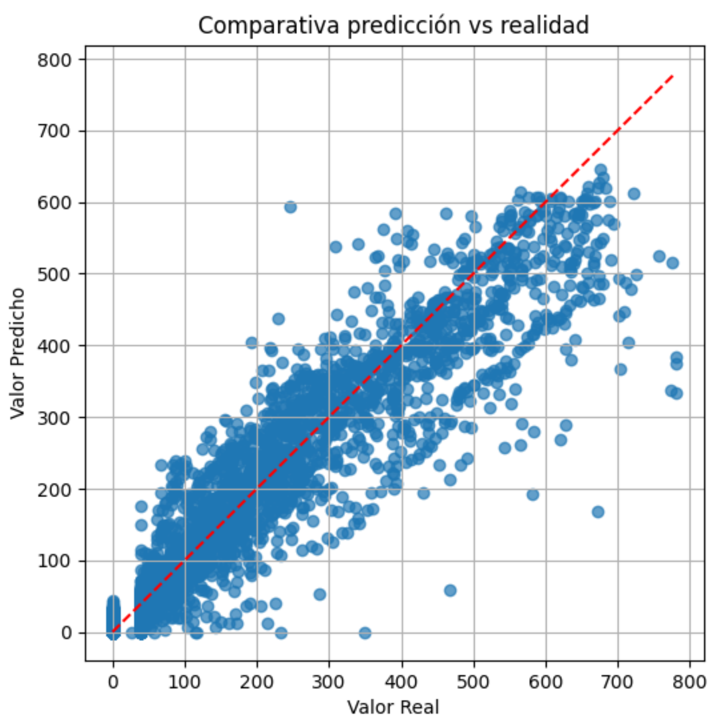

Fig. 7. Predicted vs. real values for the full test set.

5 Conclusions and future steps

The development of a prediction model does not make sense whitout a practical purpose. As was mentioned earlier, the objective of this project was to handle the electrical control of an installation under the zero inject rule. Thanks to the model, it is possible to expect how much energy you are going to produce, and adjust the electrical demand in consequence.

This would be favorable in cases where a high demandant process has to take place, because the person responsible will have information of when is the best moment to make it happen, if it needs to be made now because a fall in the production is expected, or the opposite, if it would be better to wait because a rise will happen.

It also helps the correct administration of the batteries, knowing when to save energy and when to consume it gives a valuable advantage.

Prediction may also serve as a reference of how much the installation should be producing. If it is observed that the production is way below the predicted, the operator could check if there is dirtiness or breaks causing the fall in production. This way it would be possible to avoid maintenance chores, because of the ease to check if something is not working.

In the future, it would be interesting to syncronize the model to AEMET in real time. AEMET is the entity which provides the meteorological data daily in Spain. Doing so would allow the model to work independently, and not need someone to manually tell it the exogenous variables.

Also, implementing a graphical interface would ease the process of prediction to anyone not familiarized with this kind of models. For example, designing a UI that shows you the prediction for the rest of the day and the following day, and allows you to view the prediction for a day in the future in particular.

Disclosure of Interests. The authors have no competing interests to declare that are relevant to the content of this article.

References

1. Ministerio para la Transición Ecológica: Real Decreto 244/2019, de 5 de abril, por el que se regulan las condiciones administrativas, técnicas y económicas del autoconsumo de energía eléctrica, vol. BOE-A-2019-5089 (2019) 35674–35719. Accedido: 23 de febrero de 2026. [En línea]. Disponible en: https://www.boe.es/eli/es/rd/2019/04/05/244
2. Castán-Lascorz, M. A., Mselle, B. D., Zabaleta, L. L., León, M., Arroyo, J.: Energy Flexibility Optimization in Industry: A Hybrid Approach with Synthetic Data Evaluation. In: 2025 IEEE International Conference on Engineering, Technology, and Innovation (ICE/ITMC), Valencia, Spain. IEEE (Jun. 2025) 1–10. https://doi.org/10.1109/ICE/ITMC65658.2025.11106527
3. Interquartile Range (IQR): What it is and How to Find it. Statistics How To. Accedido: 1 de diciembre de 2025. [En línea]. Disponible en: https://www.statisticshowto.com/probability-and-statistics/interquartile-range/
4. Energía solar fotovoltaica: potencia instalada en España 2010–2024. Statista. Accedido: 30 de octubre de 2025. [En línea]. Disponible en: https://es.statista.com/estadisticas/1004387/potencia-solar-fotovoltaica-instalada-enespana/
5. Modelo de Machine Learning para Interpolar Espacialmente Funciones Respuestas Impulsivas. Scribd. Accedido: 2 de enero de 2026. [En línea]. Disponible en: https://es.scribd.com/document/670410396/Modelo-de-Machine-Learning-parainterpolar-espacialmente-funciones-respuestas-impulsivas
6. Amat Rodrigo, J., Escobar Ortiz, J.: skforecast (27 de septiembre de 2023). Zenodo. https://doi.org/10.5281/ZENODO.8382788
7. Por qué y cuándo escalar los datos. Codificando Bits. Accedido: 7 de enero de 2026. [En línea]. Disponible en: https://codificandobits.com/tutorial/porquecuando-escalar-los-datos/
8. Optimización energética industrial con control predictivo. RAutomation. Accedido: 23 de febrero de 2026. [En línea]. Disponible en: https://rautomation.es/2025/09/25/optimizacion-energetica-industrial-con-controlpredictivo/
9. Audema: Situación actual de la energía fotovoltaica en España - Audema. Accedido: 3 de noviembre de 2025. [En línea]. Disponible en: https://www.audema.com/situacion-actual-de-la-energia-fotovoltaica-en-espana/
10. lafabricasolar_vmphyw: El impacto de la energía solar fotovoltaica en la economía española. La Fábrica Solar. Accedido: 4 de noviembre de 2025. [En línea]. Disponible en: https://www.lafabricasolar.com/impacto-energia-solarfotovoltaica/
11. ScienceDirect: Article S2352484725001490. Accedido: 6 de marzo de 2026. [En l'inea]. Disponible en: https://www.sciencedirect.com/science/article/pii/S2352484725001490
12. IJPEDS: Article download 24057/15126. Accedido: 6 de marzo de 2026. [En línea]. Disponible en: https://ijpeds.iaescore.com/index.php/IJPEDS/article/download/24057/15126
13. PMC: Article PMC11002275. Accedido: 6 de marzo de 2026. [En línea]. Disponible en: https://pmc.ncbi.nlm.nih.gov/articles/PMC11002275/
14. PV Magazine: Nueva base de datos abierta para estimar la producci'on fotovoltaica en España. Accedido: 6 de marzo de 2026. [En l'inea]. Disponible en: https://www.pv-magazine.es/2025/02/03/nueva-base-de-datos-abierta-para-estimar-la-produccion-fotovoltaica-en-espana/